

Architectures

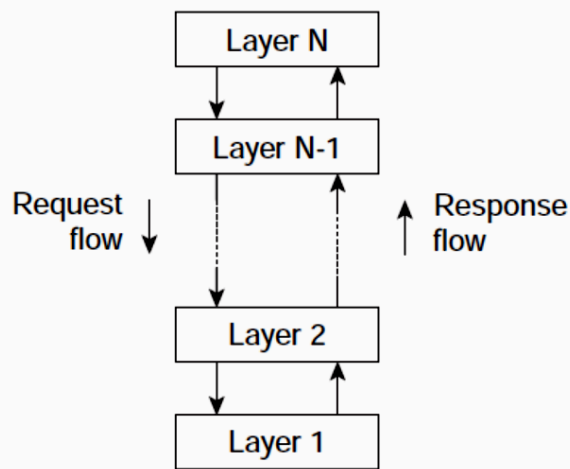
Distributed Systems [2]

Distributed System Architecture


- A *distributed system* is many cooperating computers that appear to users as a single service.
- Distributed systems are often complex pieces of software of which the components are by definition dispersed across multiple machines.
- The organization of distributed systems is mostly about the software components that constitute the system.

Architectural Styles

- Organize into logically different components, and distribute those components over the various machines.



(a)



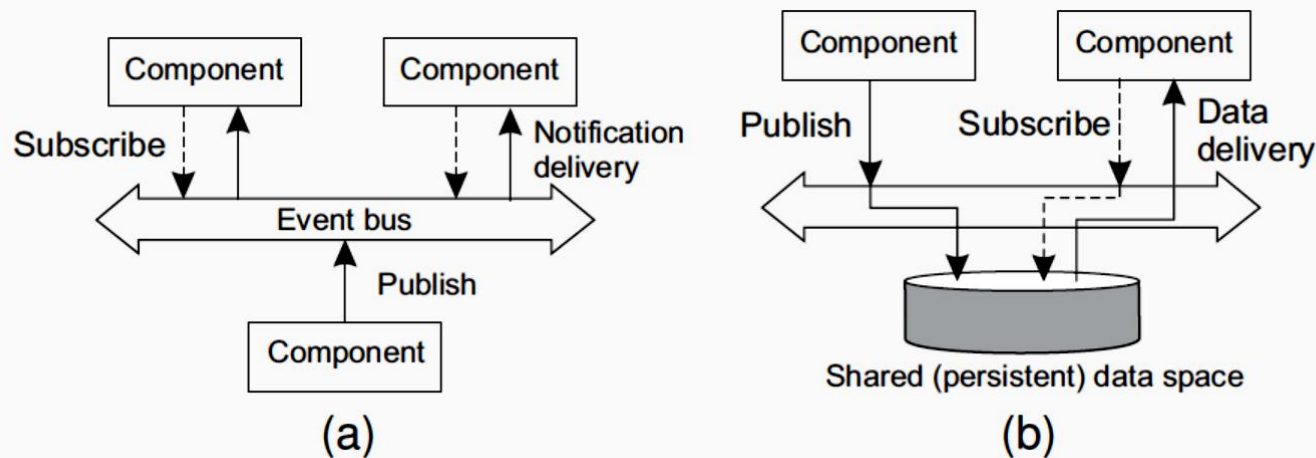
(b)

(a) Layered style is used for client-server system.

(b) Object-based style for distributed object systems.

Architectural Styles

- Decoupling processes in space (“anonymous”) and also time (“asynchronous”) has led to alternative styles.



(a) Publish/subscribe [decoupled in **space**]

(b) Shared dataspace [decoupled in **space** and **time**]


Organization

- Centralized
- Decentralized
- Hybrid

Centralized Architecture

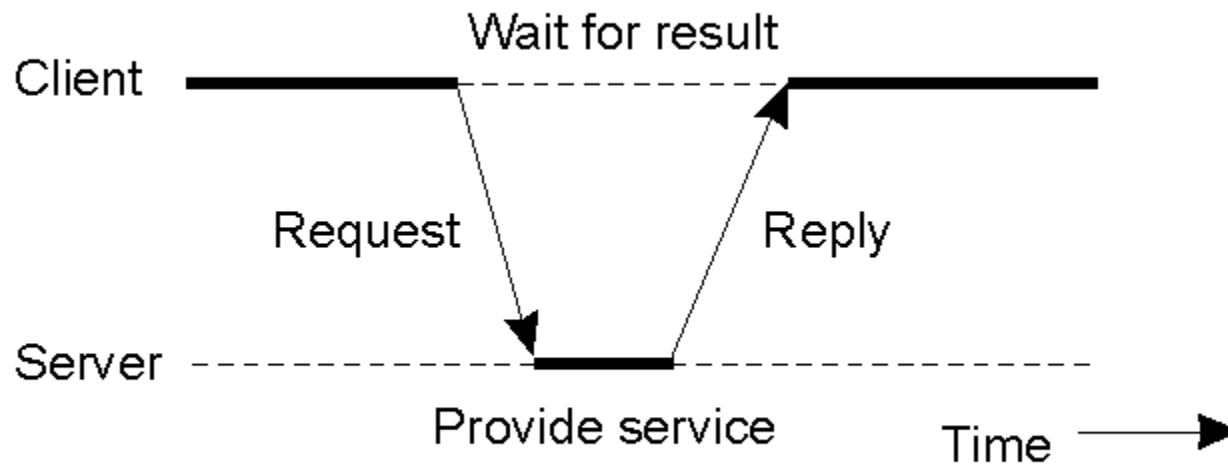
- **Basic Client–Server Model**
- Characteristics:
 - There are processes offering services (servers)
 - There are processes that use services (clients)
 - Clients and servers can be on different machines
 - Clients follow request/reply model wrt to using services

Client-Server Communication

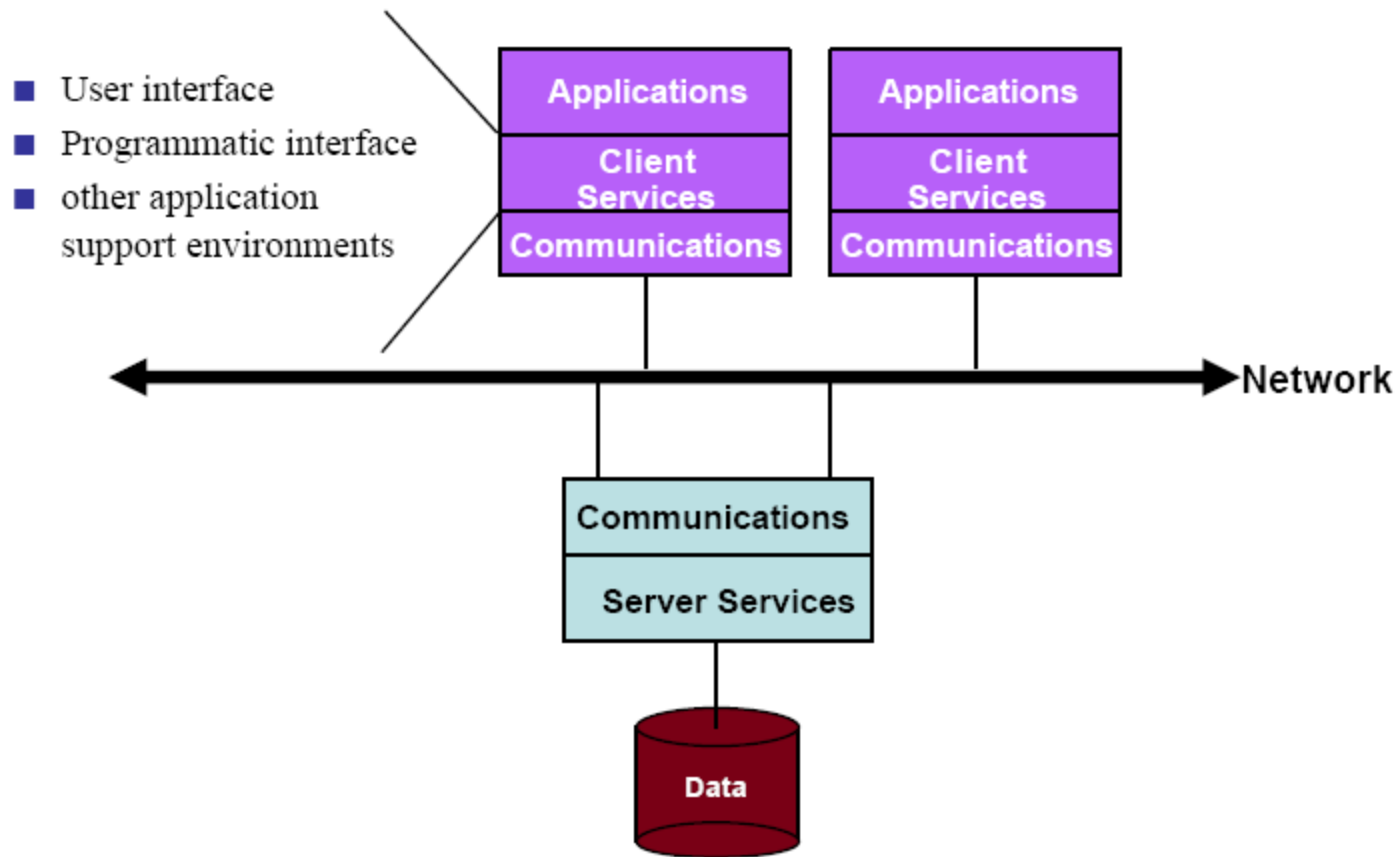


Clients and Servers, Timing

- General interaction between a client and a server.




Multiple-Client/Single Server




Problems With Multiple-Client/Single Server

- Server forms bottleneck
- Server forms single point of failure
- System scaling difficult


A service provided by multiple servers



Multiple Clients/Multiple Servers




Web proxy server



Web Applets

a) client request results in the downloading of applet code



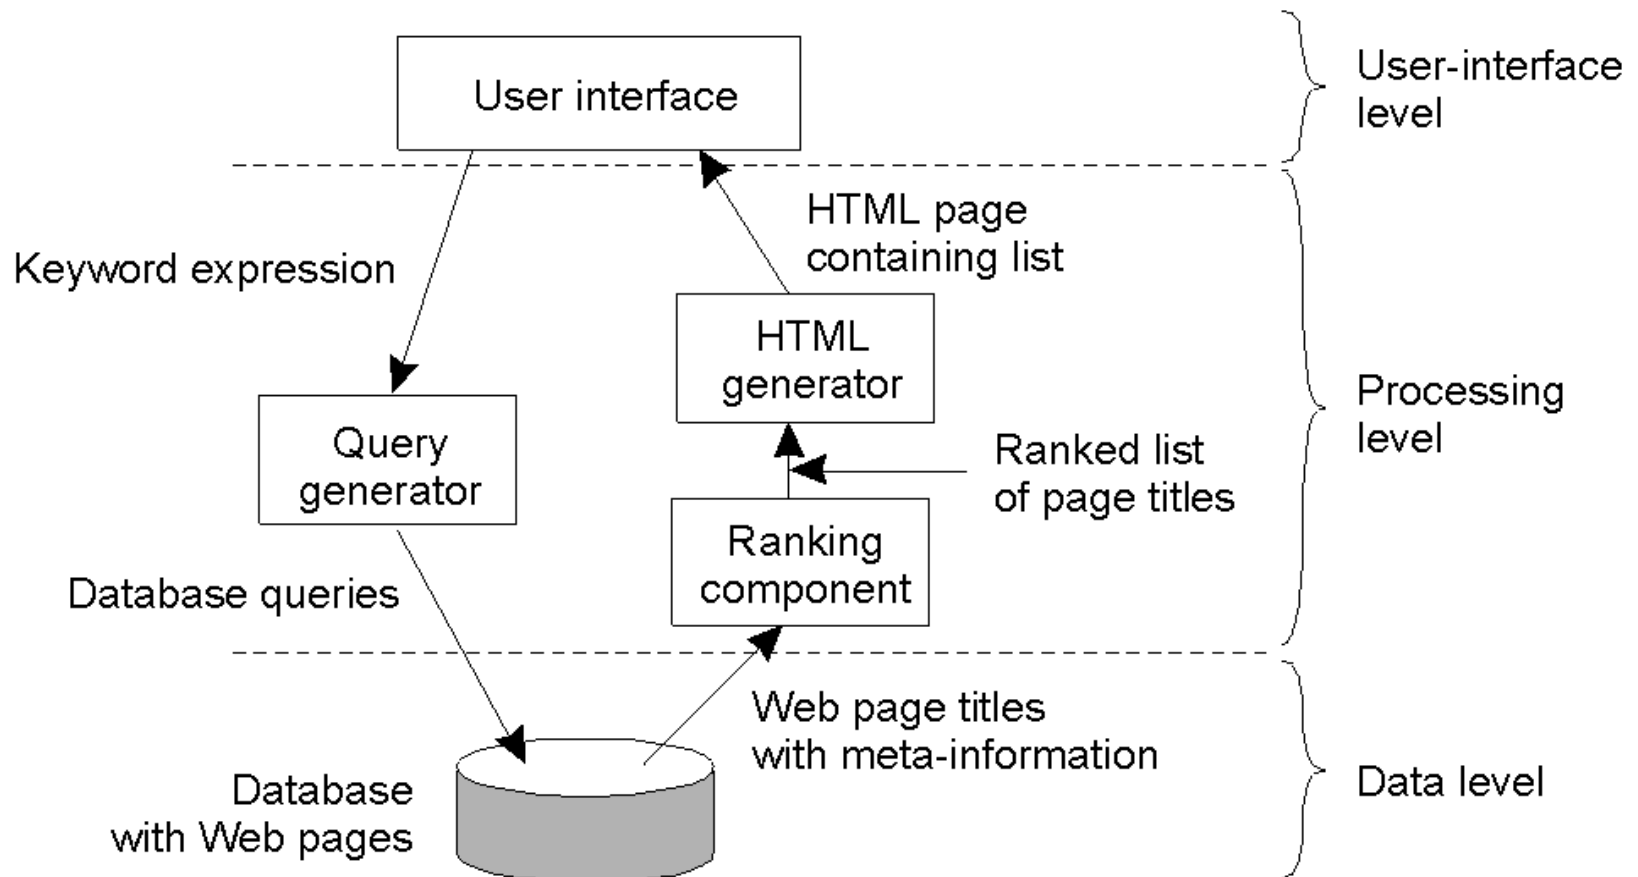
b) client interacts with the applet



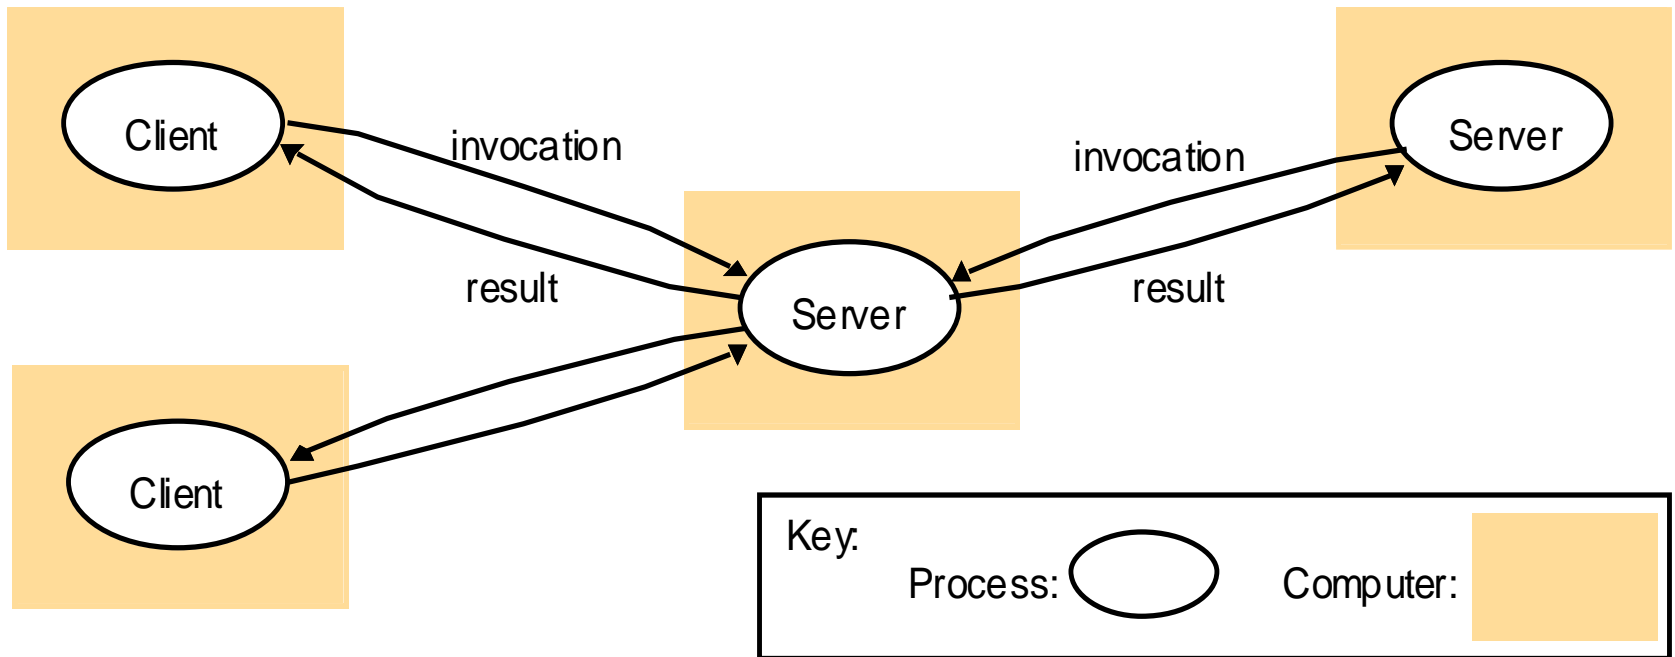
Application Layering

- **Traditional three-layered view**
 - User-interface layer contains units for an application's user interface
 - Processing layer contains the functions of an application, i.e. without specific data
 - Data layer contains the data that a client wants to manipulate through the application components
- This layering is found in many distributed information systems, using traditional database technology and accompanying applications.

Processing Level

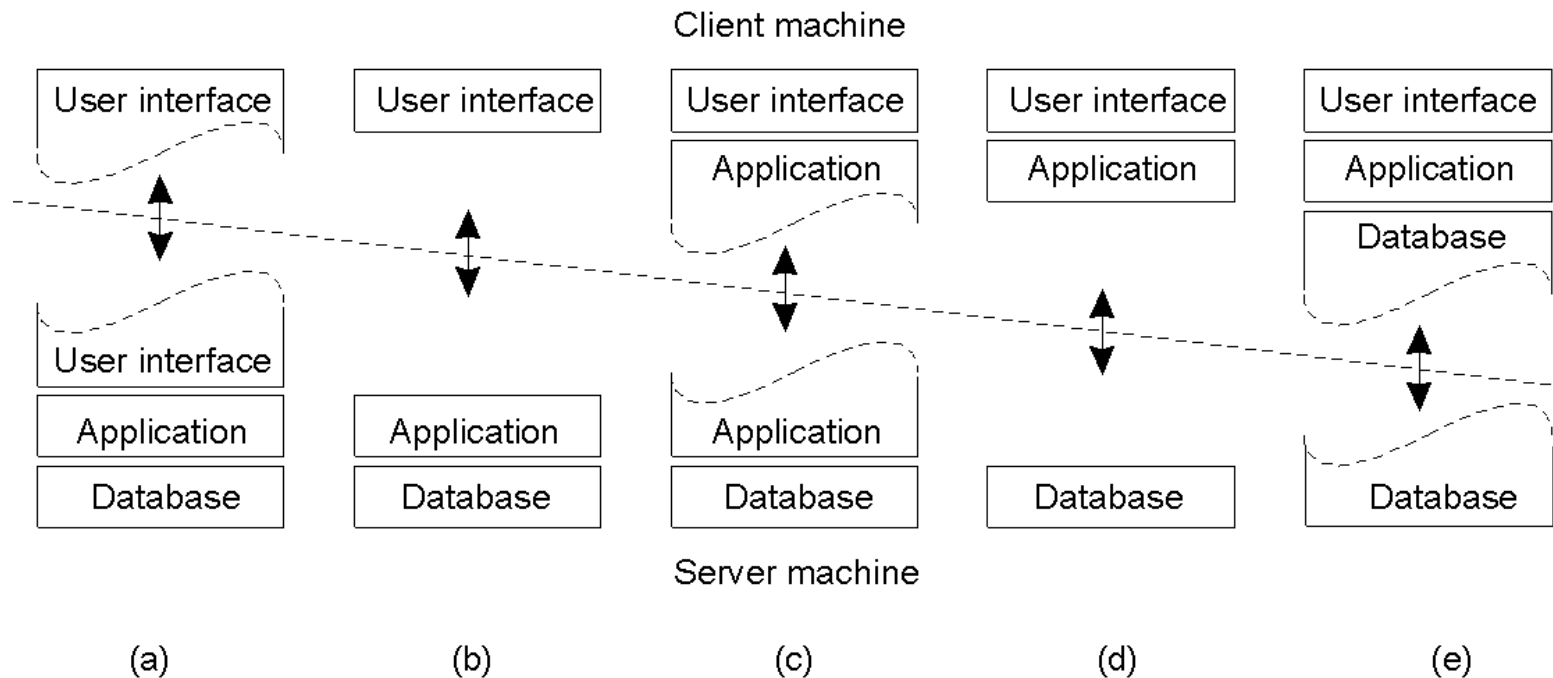


Clients Invoke Individual Servers



Multitiered Architectures (1)

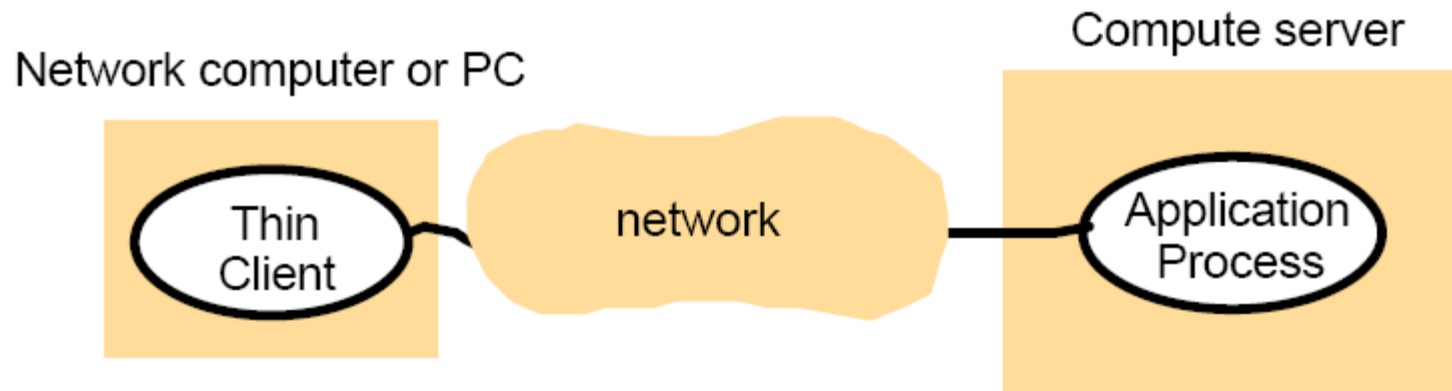
- Single-tiered: dumb terminal/mainframe configuration
- Two-tiered: client/single server configuration
- Three-tiered: each layer on separate machine



“Thin” Clients

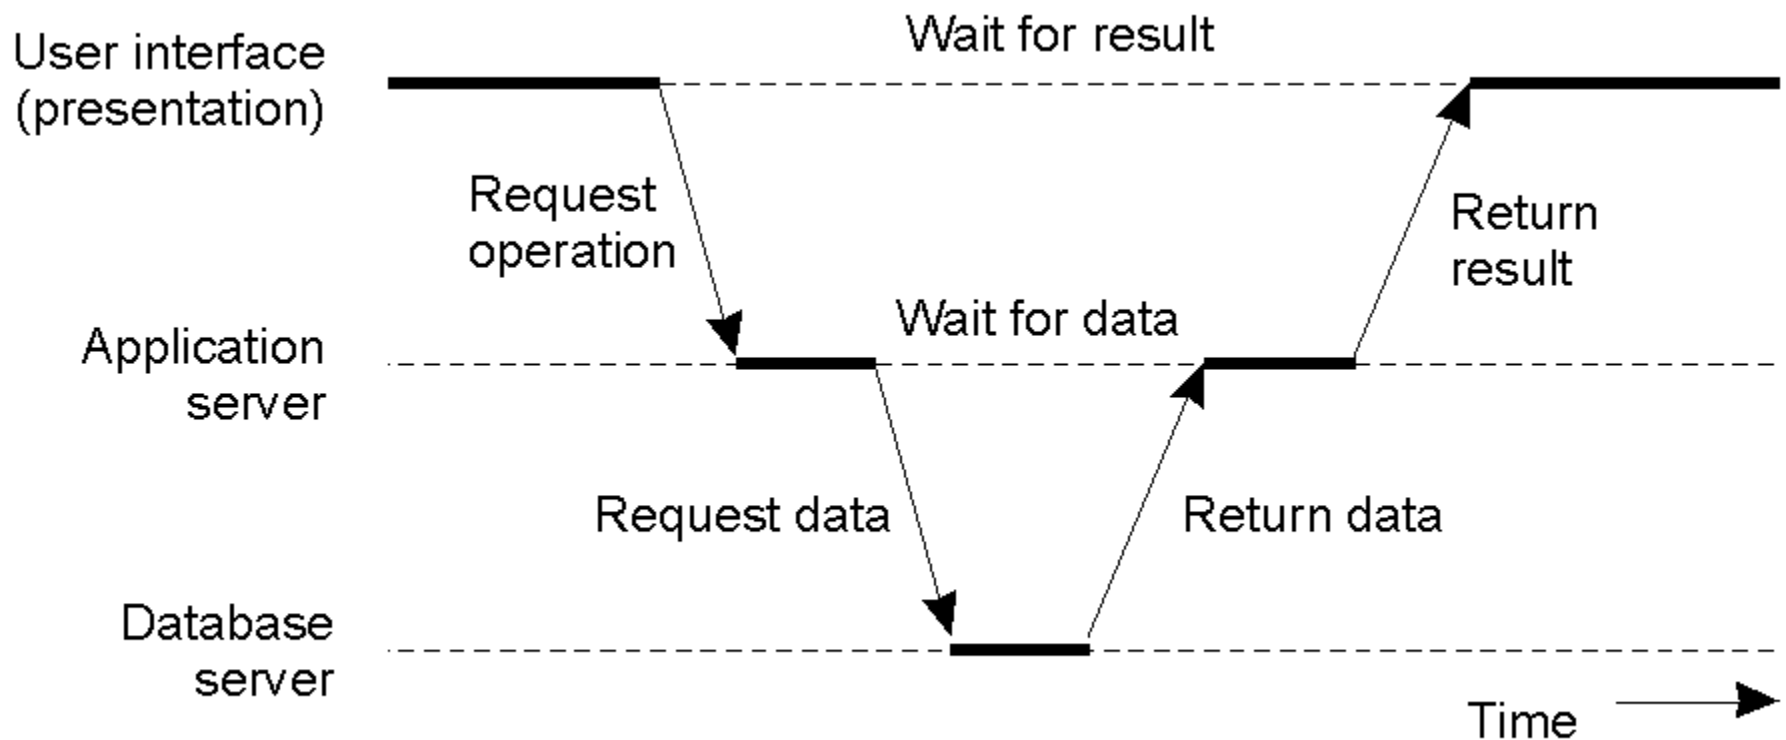
■ Thin Clients and Compute Servers

- Executing graphical user interface on local computer while application executes on compute server
- Example: X11 server (run on the application client side)
- In reality: Palm Pilots, Mobile phones



Multitiered Architectures (2)


- An example of a server acting as a client.



Decentralized Architecture



- Structured P2P: nodes are organized following a specific distributed data structure
- Unstructured P2P: nodes have randomly selected neighbors
- Hybrid P2P: some nodes are appointed special functions in a well-organized fashion
- In virtually all cases, we are dealing with overlay networks: data is routed over connections setup between the nodes (cf. application-level multicasting)

A Distributed Application based on Peer Processes



Structured P2P Systems

- Organize the nodes in a structured overlay network such as a logical ring, or a hypercube, and make specific nodes responsible for services based only on their ID.




Unstructured P2P Systems

- Many unstructured P2P systems are organized as a random overlay: two nodes are linked with probability p .
- We can no longer look up information deterministically, but will have to resort to searching:
 - Flooding: node u sends a lookup query to all of its neighbors. A neighbor responds, or forwards (floods) the request. There are many variations:
 - Limited flooding (maximal number of forwarding)
 - Probabilistic flooding (flood only with a certain probability).
 - Random walk: Randomly select a neighbor v . If v has the answer, it replies, otherwise v randomly selects one of its neighbors.
Variation: parallel random walk. Works well with replicated data.


Superpeers

- Sometimes it helps to select a few nodes to do specific work: superpeer.
 - Peers maintaining an index (for search)
 - Peers monitoring the state of the network
 - Peers being able to setup connections



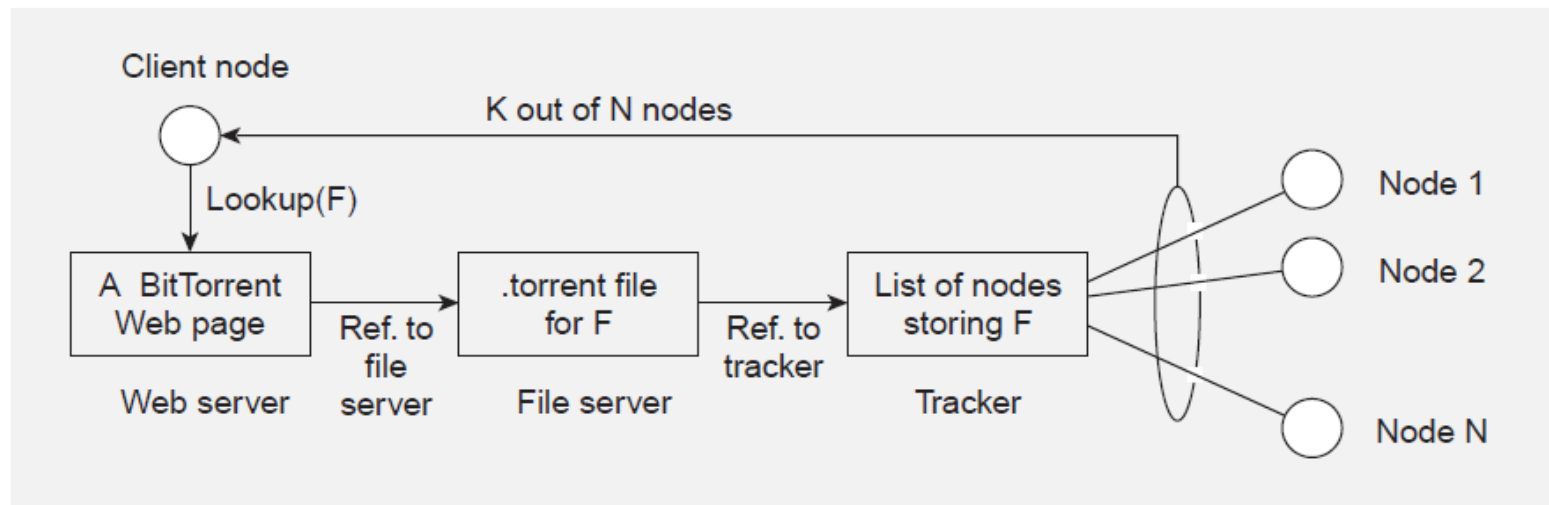
Hybrid Architectures

- Client-server combined with P2P
- Edge-server architectures, which are often used for Content Delivery Networks.



BitTorrent


- Once a node has identified where to download a file from, it joins a swarm of downloaders who in parallel get file chunks from the source, but also distribute these chunks amongst each other.



Middleware

- In many cases, distributed systems/applications are developed according to a specific architectural style. The chosen style may not be optimal in all cases → need to (dynamically) adapt the behavior of the middleware.
- **Interceptors**
 - Intercept the usual flow of control when invoking a remote object.

A distributed system organized as middleware



Self-managing Distributed Systems

- Distinction between system and software architectures blurs when automatic adaptively needs to be taken into account:
 - Self-configuration
 - Self-managing
 - Self-healing
 - Self-optimizing
 - ...

There is a lot of hype going on in this field of **autonomic computing**.

Feedback Control Model

- In many cases, self-* systems are organized as a feedback control system.

