

# Processes

Distributed Systems [3]

# What is a Process?

- Process: An **execution stream** in the context of a **process state**
- **Execution stream**
  - Stream of executing instructions
  - Running piece of code
  - Sequential sequence of instructions
  - “thread of control”
- **Process state**
  - Everything that the running code can affect or be affected by

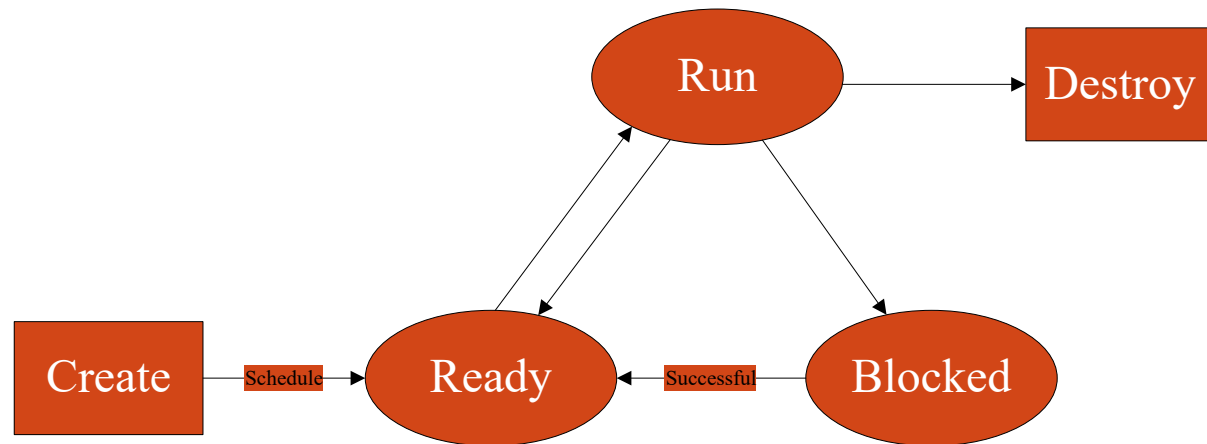
# Processes vs. Programs

- A process is different than a program
  - Program: Static code and static data
  - Process: Dynamic instance of code and data
- No one-to-one mapping between programs and processes
  - Can have multiple processes of the same program  
Example: many users can run “ls” at the same time
  - One program can invoke multiple processes  
Example: “make” runs many processes to accomplish its work

VNF vs. VNFI

# Processes

- Each process is in one of three modes:
  - Running: On the CPU (only one on a uniprocessor)
  - Ready: Waiting for the CPU
  - Blocked (or asleep): Waiting for I/O or synchronization to complete



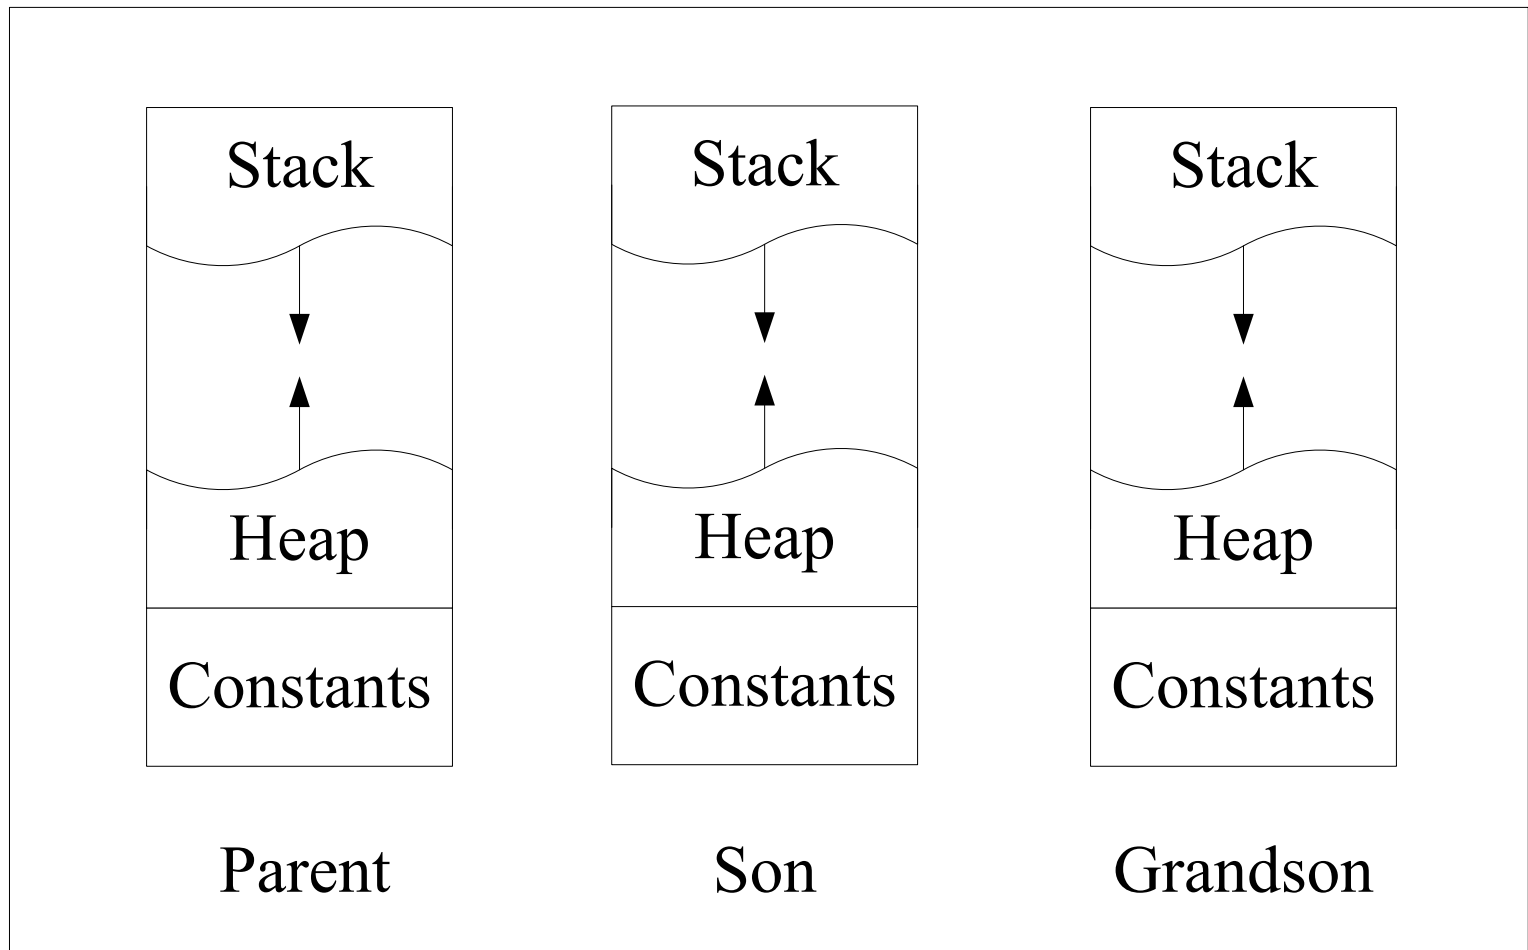
# Create Processes

- `/*Unix process_creation.c */`

```
#include <stdlib.h>
```

```
Main () {  
    int pid;  
    if ((pid = fork ()) != 0) {  
        printf ("Father\n"); wait(0);  
    }  
    else if ((pid = fork ()) != 0){  
        printf ("Son\n"); wait (0);  
    }  
    else printf ("Grandson\n");  
    exit (0);  
}
```

# Memory Created in UNIX Processes



# Low Performance

- Resource management
  - When creating a new process, assign address space, copy data
- Scheduling
  - Context switch
    - Process Context: CPU context and storage context
- Cooperation
  - IPC, interprocess communication
  - Shared memory

# Introduction to Threads

- Thread: A minimal software processor in whose context a series of instructions can be executed. Saving a thread context implies stopping the current execution and saving all the data needed to continue the execution at a later stage.



# Context Switching

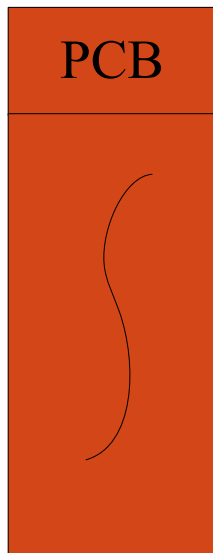
- Processor context: The minimal collection of values stored in the registers of a processor used for the execution of a series of instructions (e.g., stack pointer, addressing registers, program counter).
- Thread context: The minimal collection of values stored in registers and memory, used for the execution of a series of instructions (i.e., processor context, state).
- Process context: The minimal collection of values stored in registers and memory, used for the execution of a thread (i.e., thread context, but now also at least MMU register values).

# Context Switching


- Threads share the same address space. Thread context switching can be done entirely independent of the operating system.
- Process switching is generally more expensive as it involves getting the OS in the loop, i.e., trapping to the kernel.
- Creating and destroying threads is much cheaper than doing so for processes.

# Thread and Process

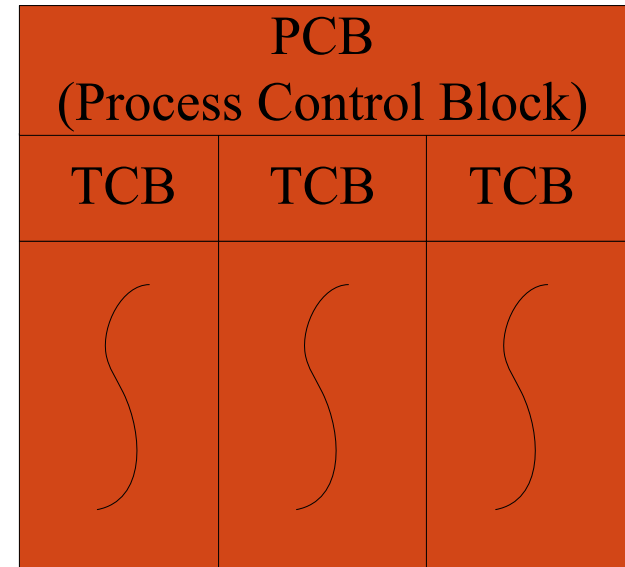
Single Thread  
Process



Multi Thread  
Process



Multi Thread  
Process



Tread System

Operation System


# User-Level Thread

- All the threads are created in user processes' address spaces.
- Advantage
  - All operations can be completely handled within a single process  $\Rightarrow$  implementations can be extremely efficient.
- Disadvantage
  - Difficult to get the support from OS, block
    - All services provided by the kernel are done on behalf of the process in which a thread resides  $\Rightarrow$  if the kernel decides to block a thread, the entire process will be blocked.

# Threads and Operating Systems

- Have the kernel contain the implementation of a thread package. This means that all operations return as system calls:
  - Operations that block a thread are no longer a problem: the kernel schedules another available thread within the same process.
  - Handling external events is simple: the kernel (which catches all events) schedules the thread associated with the event.
  - The problem is (or used to be) the loss of efficiency due to the fact that each thread operation requires a trap to the kernel.

# Thread Implementation



# Threads and Operating Systems

- Conclusion – but
  - Try to mix user-level and kernel-level threads into a single concept, however, performance gain has not turned out to outweigh the increased complexity.

# Threads and Distributed Systems

- Hiding network latencies:
  - Web browser scans an incoming HTML page, and finds that more files need to be fetched.
  - Each file is fetched by a separate thread, each doing a (blocking) HTTP request.
  - As files come in, the browser displays them.
- Multiple request-response calls to other machines(RPC)
  - A client does several calls at the same time, each one by a different thread.
  - It then waits until all results have been returned.
  - Note: if calls are to different servers, we may have a linear speed-up.

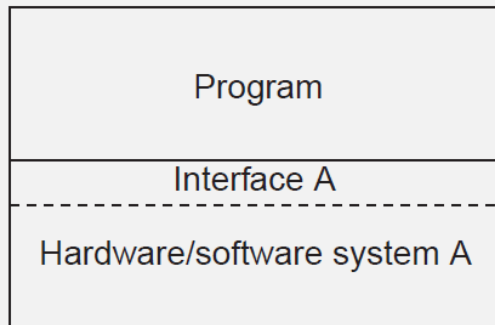


# Threads and Distributed Systems

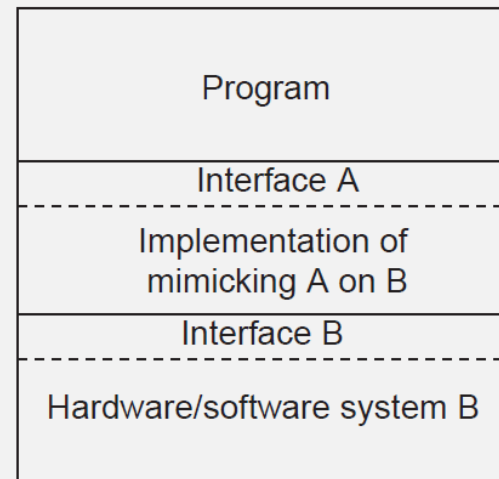
- Improve performance:
  - Starting a thread is much cheaper than starting a new process.
  - Having a single-threaded server prohibits simple scale-up to a multiprocessor system.
  - As with clients: hide network latency by reacting to next request while previous one is being replied.
- Better structure
  - Most servers have high I/O demands. Using simple, well-understood blocking calls simplifies the overall structure.
  - Multithreaded programs tend to be smaller and easier to understand due to simplified flow of control.

# Virtualization

- Virtualization is becoming increasingly important:
  - Hardware changes faster than software
  - Ease of portability and code migration
  - Isolation of failing or attacked components



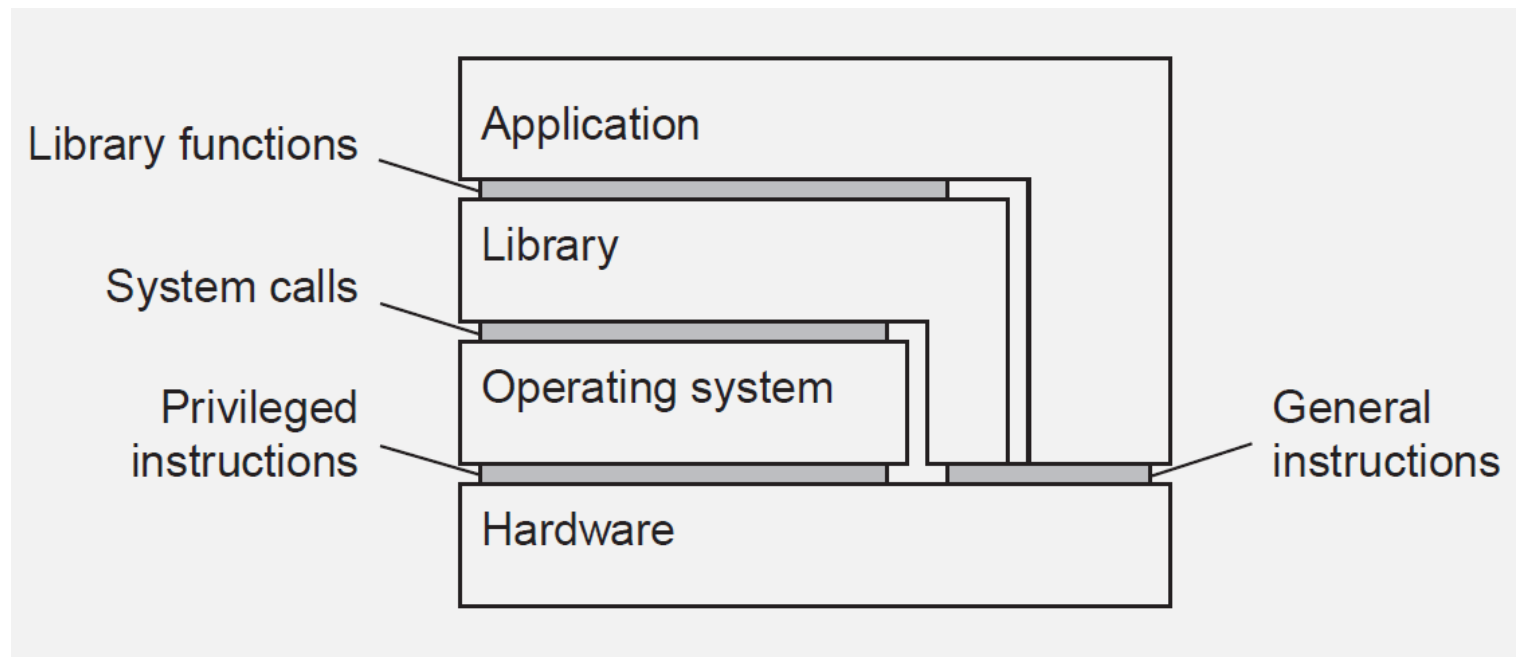
(a)



(b)

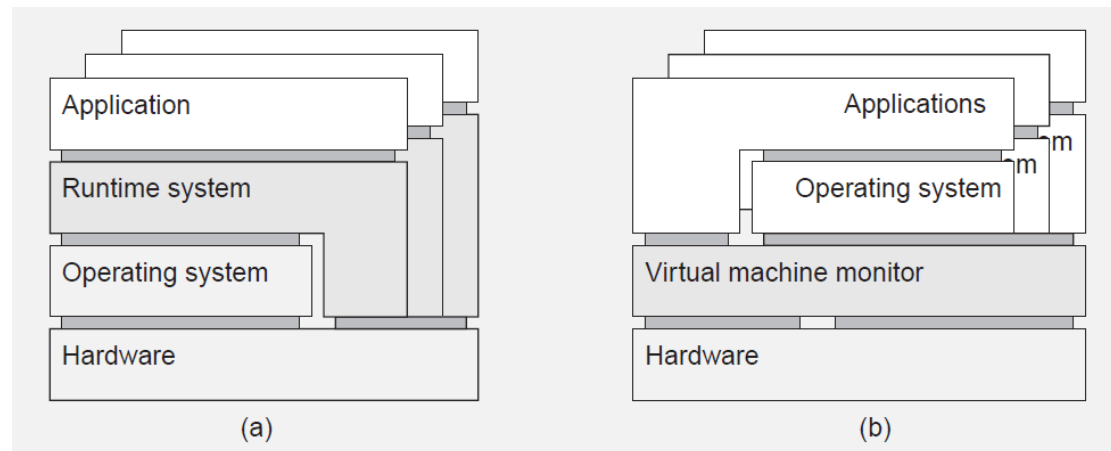
# Architecture of VMs

- Virtualization can take place at very different levels, strongly depending on the interfaces as offered by various systems components:



# Process VMs versus VM Monitors

- Process VM: A program is compiled to intermediate (portable) code, which is then executed by a runtime system (Example: Java VM).
- VM Monitor: A separate software layer mimics the instruction set of hardware  $\Rightarrow$  a complete operating system and its applications can be supported (Example: VMware, KVM, XEN).

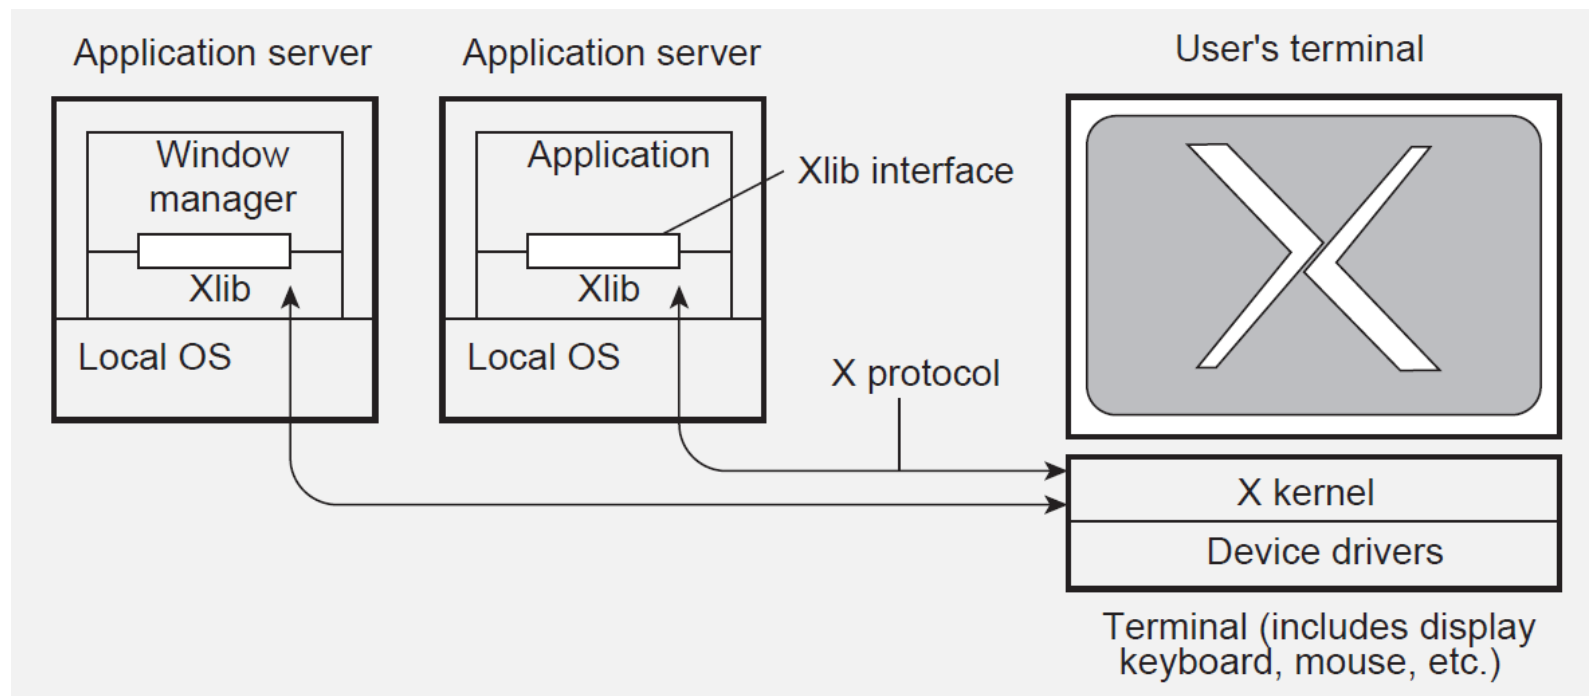


# VM Monitors on operating systems

- We're seeing VMMs run on top of existing operating systems
  - Perform binary translation: while executing an application or operating system, translate instructions to that of the underlying machine.
  - Distinguish sensitive instructions: traps to the original kernel (think of system calls, or privileged instructions).
  - Sensitive instructions are replaced with calls to the VMM.

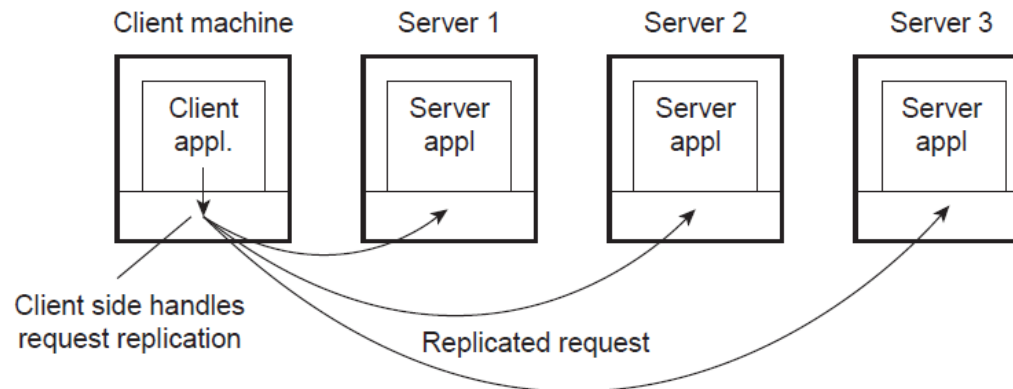
# Clients: User Interfaces

- A major part of client-side software is focused on (graphical) user interfaces.



# Client-Side Software

- Generally tailored for distribution transparency
  - access transparency: client-side stubs for RPCs
  - location/migration transparency: let client-side software keep track of actual location
  - failure transparency: can often be placed only at client (we're trying to mask server and communication failures).
  - replication transparency: multiple invocations handled by client stub



# Multithreaded Clients

- Example, web browser

A web document  $\supset$  plain text, a collection of images.

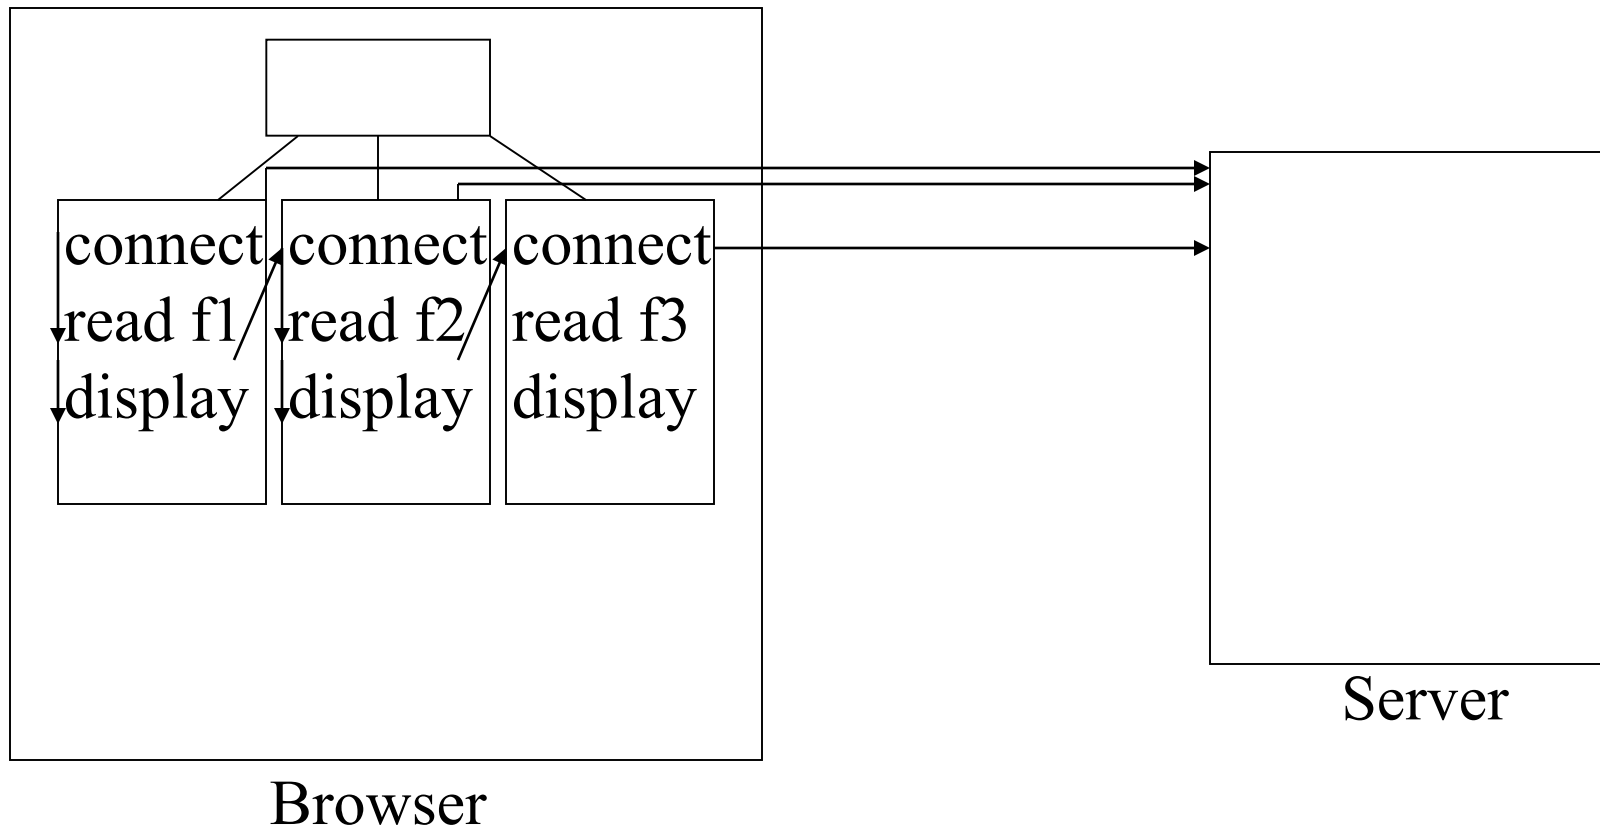
To fetch a HTML file:

`connect` server, `read` a file1, `display`

`connect` server, `read` a file2, `display`

... ..






# Servers: General organization

- Basic model:
  - A server is a process that waits for incoming service requests at a specific transport address. In practice, there is a one-to-one mapping between a port and a service.

ftp-data	20	File Transfer [Default Data]
ftp	21	File Transfer [Control]
telnet	23	Telnet
	24	any private mail system
smtp	25	Simple Mail Transfer
login	49	Login Host Protocol
sunrpc	111	SUN RPC (portmapper)
courier	530	Xerox RPC

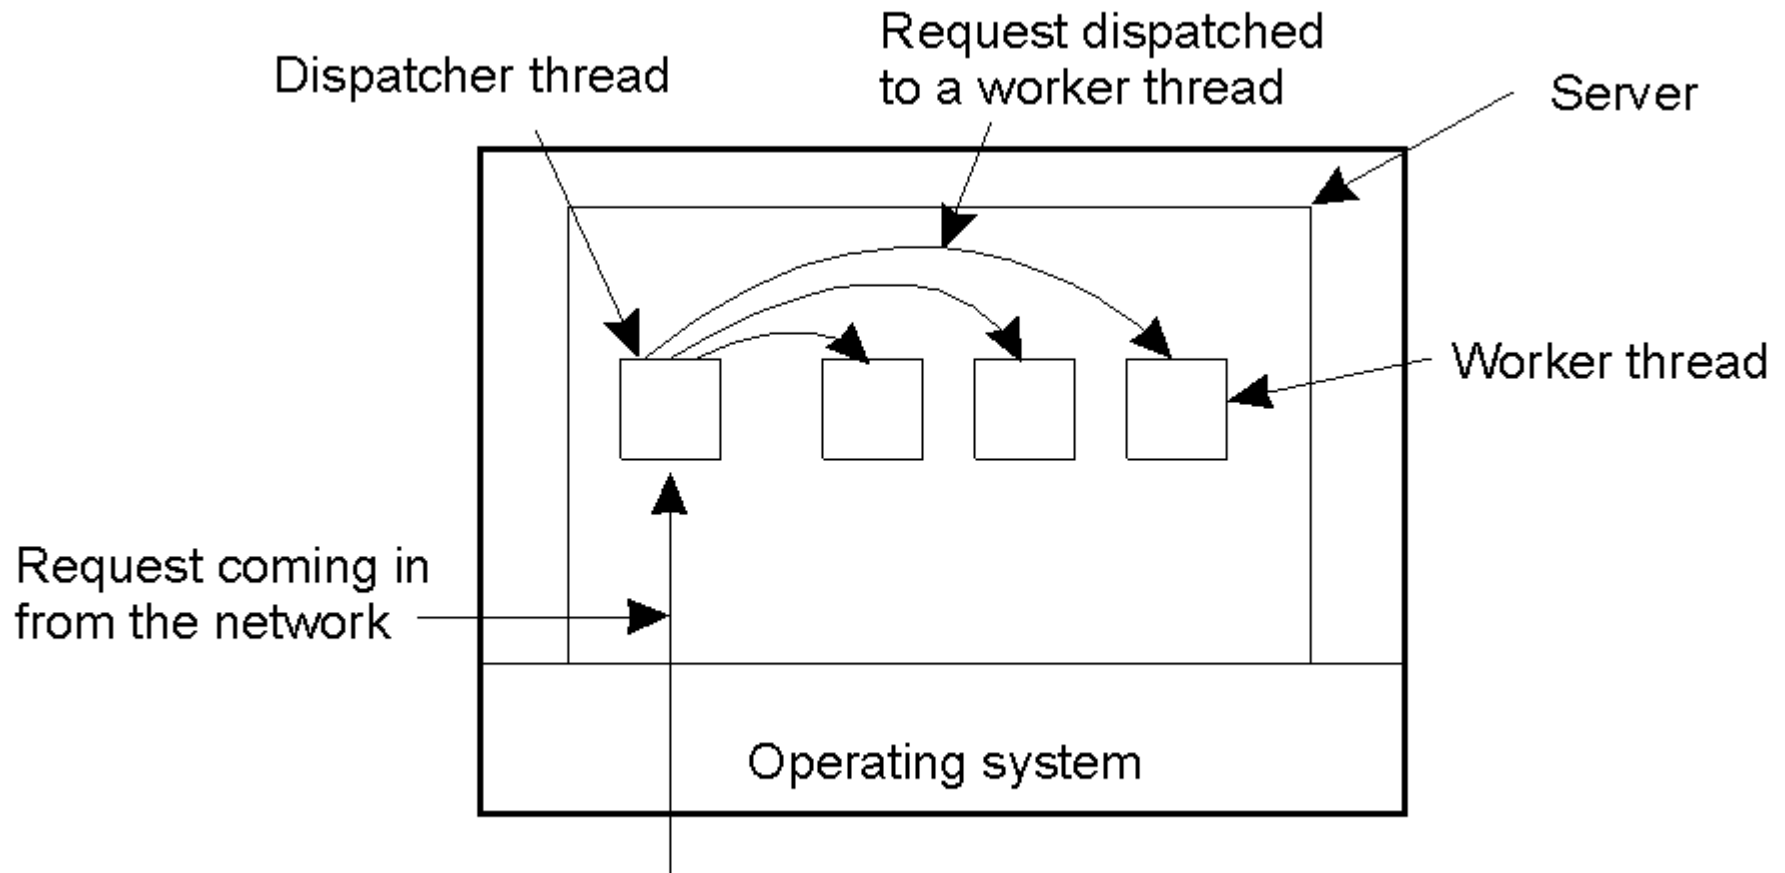
# Multithreaded Servers

- Example, file server



# Multithreaded Servers

- A multithreaded server organized in a dispatcher/worker model.



# Servers: General organization

- Type of servers:
  - Superservers: Servers that listen to several ports, i.e., provide several independent services. In practice, when a service request comes in, they start a subprocess to handle the request (UNIX inetd)
  - Iterative vs. concurrent servers: Iterative servers can handle only one client at a time, in contrast to concurrent servers

# Out-of-band communication

- Is it possible to interrupt a server once it has accepted (or is in the process of accepting) a service request?
- Use a separate port for urgent data:
  - Server has a separate thread/process for urgent messages
  - Urgent message comes in  $\Rightarrow$  associated request is put on hold
  - Note: we require OS supports priority-based scheduling
- Use out-of-band communication facilities of the transport layer:
  - Example: TCP allows for urgent messages in same connection
  - Urgent messages can be caught using OS signaling techniques

# Stateless

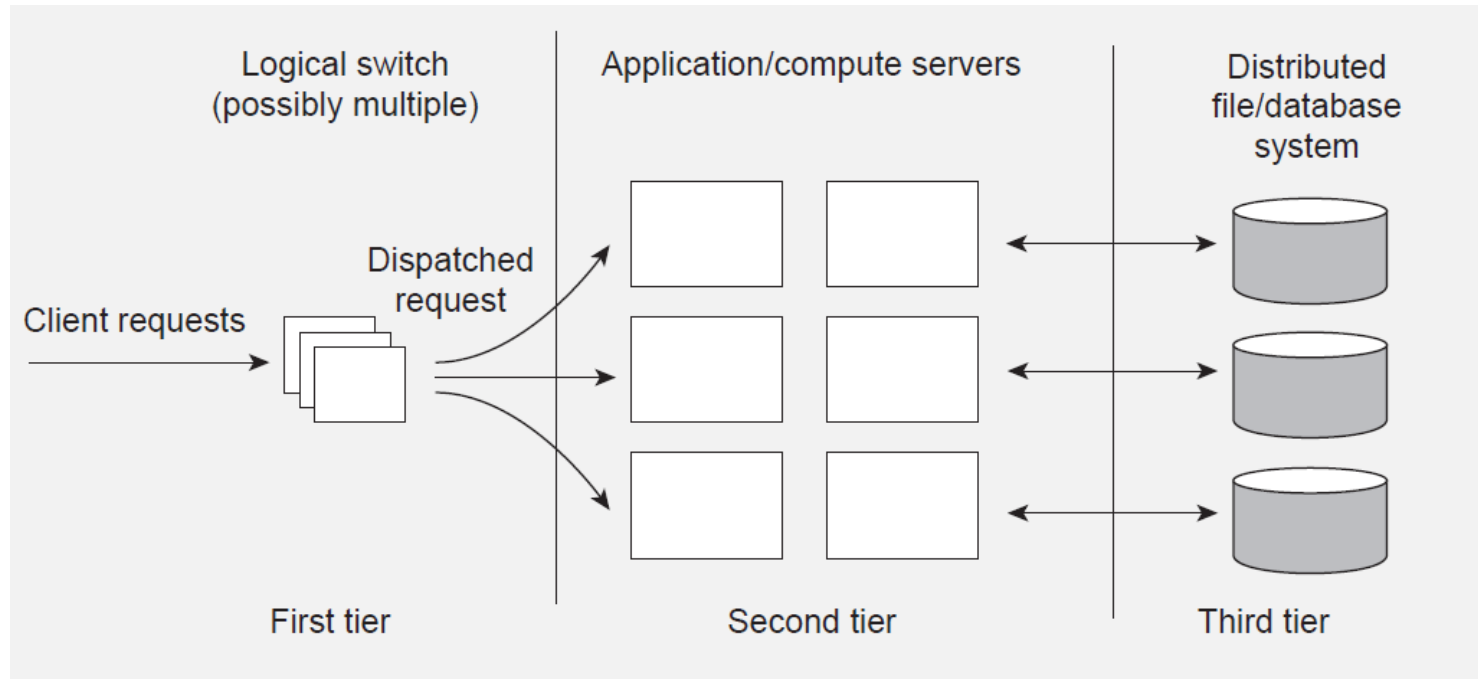
- Never keep accurate information about the status of a client after having handled a request:
  - Don't record whether a file has been opened (simply close it again after access)
  - Don't promise to invalidate a client's cache
  - Don't keep track of your clients
- Consequences
  - Clients and servers are completely independent
  - State inconsistencies due to client or server crashes are reduced
  - Possible loss of performance because, e.g., a server cannot anticipate client behavior (think of prefetching file blocks)

# Stateful

- Keeps track of the status of its clients:
  - Record that a file has been opened, so that prefetching can be done
  - Knows which data a client has cached, and allows clients to keep local copies of shared data
- The performance of stateful servers can be extremely high, provided clients are allowed to keep local copies. As it turns out, reliability is not a major problem.



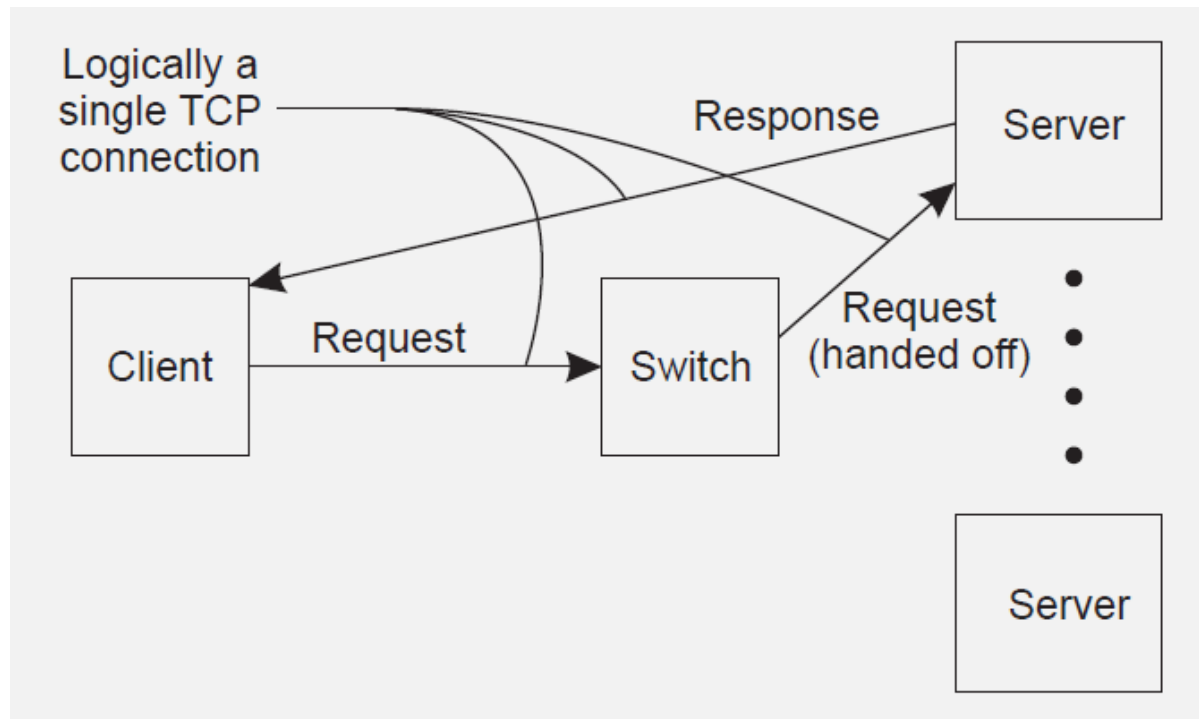
# Server clusters: three different tiers



- Crucial element
  - The first tier is generally responsible for passing requests to an appropriate server.

# Request Handling

- Having the first tier handle all communication from/to the cluster may lead to a bottleneck.
- Various, but one popular one is TCP-handoff



# Code Migration

- Approaches to code migration
- Migration and local resources
- Migration in heterogeneous systems

# Strong and weak mobility

- Code segment: contains the actual code
- Data segment: contains the state
- Execution state: contains context of thread executing the object's code

# Strong and weak mobility

- Move only code and data segment (and reboot execution):
  - Relatively simple, especially if code is portable
  - Distinguish code shipping (push) from code fetching (pull)
- Move component, including execution state
  - Migration: move entire object from one machine to the other
  - Cloning: start a clone, and set it in the same execution state

# Managing local resources

- An object uses local resources that may or may not be available at the target site
- Resource types
  - Fixed: the resource cannot be migrated, such as local hardware
  - Fastened: the resource can, in principle, be migrated but only at high cost
  - Unattached: the resource can easily be moved along with the object (e.g. a cache)

# Managing local resources

- Object-to-resource binding
  - By identifier: the object requires a specific instance of a resource (e.g. a specific database)
  - By value: the object requires the value of a resource (e.g. the set of cache entries)
  - By type: the object requires that only a type of resource is available (e.g. a color monitor)

# Managing local resources (2/2)

	Unattached	Fastened	Fixed
<b>ID</b>	MV (or GR)	GR (or MV)	GR
<b>Value</b>	CP (or MV, GR)	GR (or CP)	GR
<b>Type</b>	RB (or MV, GR)	RB (or GR, CP)	RB (or GR)

*GR = Establish global systemwide reference*

*MV = Move the resource*

*CP = Copy the value of the resource*

*RB = Re-bind to a locally available resource*



# Migration in heterogeneous systems

- Main problem
  - The target machine may not be suitable to execute the migrated code
  - The definition of process/thread/processor context is highly dependent on local hardware, operating system and runtime system
- Make use of an abstract machine that is implemented on different platforms
  - Interpreted languages, effectively having their own VM
  - Virtual VM (as discussed previously)